
Final Report of the Stock Project

Rui Li

June 15th 2023

Data Preparation

Data Perparation

- WRDS
- <https://feng-cityuhk.github.io/EquityCharacteristics/>
generate characteristics to predict stock return

Avoid data leakage

Predicting the next month's stock return using data in current months.

- `data['year_month'] = pd.to_datetime(data['date']).dt.to_period('M')`
`data['ret_fut'] = data.groupby('permno')['ret'].shift(-1)`
`data = data.dropna(axis=0, subset=['ret_fut'])` #drop column with na
- Data preprocessing within each rolling window

Architecture

```

for test_date in test_dates:
    # training set
    X_ptrain, y_ptrain = data[data.date < test_date]
    # test set
    X_test, y_test = data[data.date == test_date]

    # Hyperparameter tuning
    if test_dates.index == 0 or test_dates.index % 12 == 11:

        # training set for fine tuning
        X_train, y_train = data[data.date < test_date - 2 month]
        # 2 months span validation set
        X_valid, y_valid = data[(data.date >= test_date - 2 month)
                                &(data.date < test_date)]

        """
        Use optuna to do Hyperparameter fine tuning
        """

    # train model
    model = Model(**best_params)
    model.fit(X_ptrain, y_ptrain)
    y_pred = model.predict(X_test)

```

Rolling windows:

- A monthly expand training set
- Initialize with time span of 32 months
- Data preprocessing within each rolling
- Fine tuning per year
- Creating validation set with the last two months each year

Data :

- Data for training model: (X_ptrain , y_ptrain)
- Data for fine tuning: (X_train , y_train) , (X_valid , y_valid)
- Data for prediction and calculate R^2 : (X_test , y_test)

Model:

- OLS , lasso
- (No justifications for fine-tuning)
- RF , Xgboost , Lightgbm , pca + Lightgbm
- NN3

```

for test_date in test_dates:
    # training set
    X_ptrain, y_ptrain = data[data.date < test_date]
    #test set
    X_test, y_test = data[data.date == test_date]

    #drop categorical variables
    s = (data.dtypes == 'int64')
    object_cols = list(s[s].index)
    data = data.drop(object_cols, axis=1)

    #standardize by column

    scaler = StandardScaler()

    for col in object_cols:
        scaler.fit(X_test[col].values.reshape(-1,1))

        X_test[col] =
        scaler.transform(X_test[col].values.reshape(-1,1))

        scaler.fit(X_ptrain[col].values.reshape(-1,1))

        X_ptrain[col] =
        scaler.transform(X_ptrain[col].values.reshape(-1,1))

```

Data Preprocessing within rolling

- standardize by column within each rolling window
- drop categorical variables for OLS, Lasso and NN3
- drop future return to avoid data leakage


Method and Result

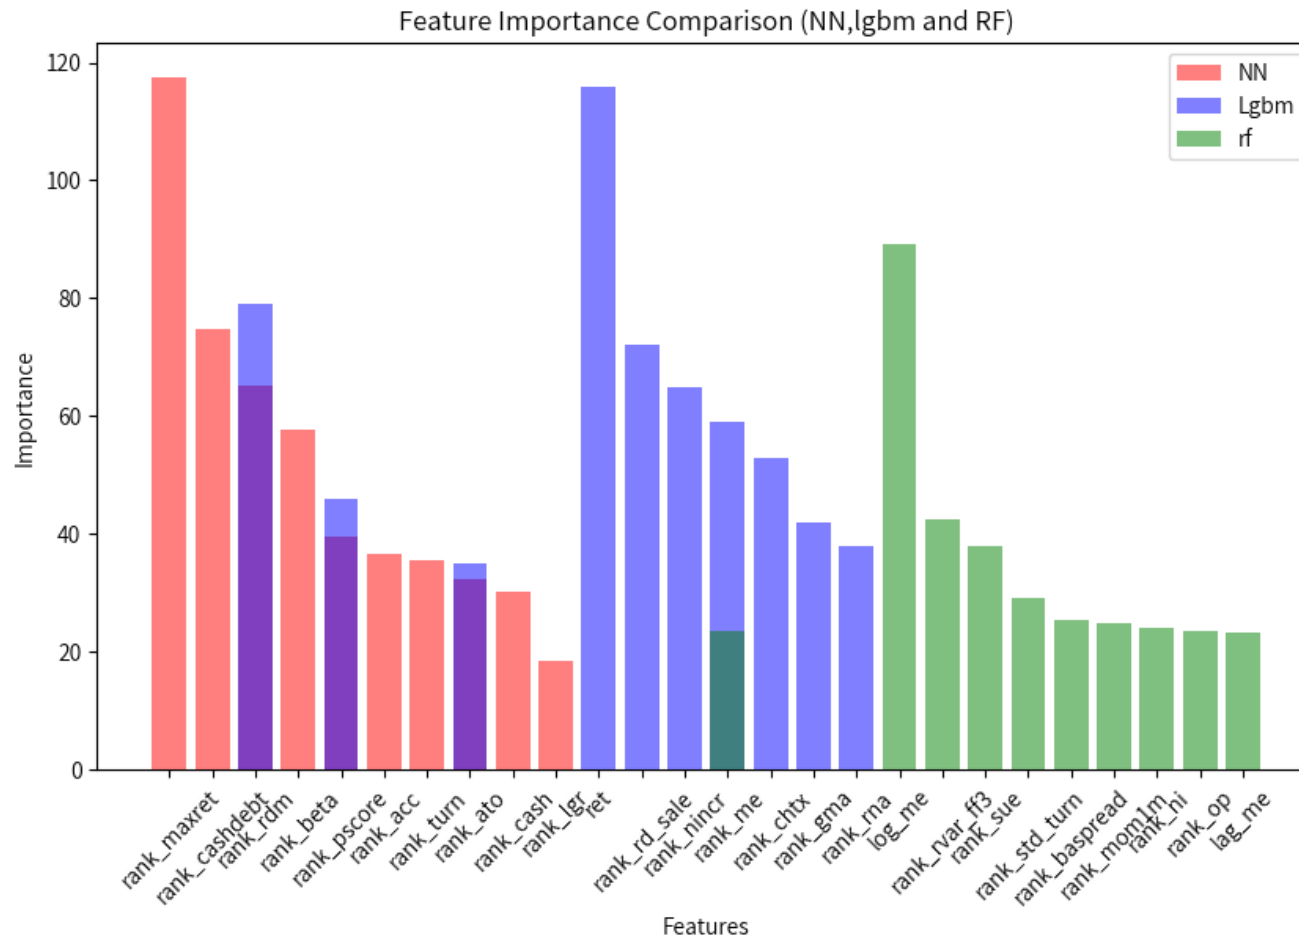
Models and Hyperparameters

Random Forest	GBRT	XGBoost	LightGBM	NN3
<code>max_depth = [1, 6]</code> <code>n_estimators = 300</code> <code>max_features</code> <code>= {3, 5, 10, 20, 30, 50,...}</code>	<code>max_depth = [1, 2]</code> <code>learning_rate = {0.001,0.1}</code> <code>n_estimators = [1,1000]</code>			<code>l1 = (1e-5 , 1e-3)</code> <code>learning_rate = {1e-3, 1e-2}</code> <code>batch_size = 10000</code> <code>epochs = 100</code> <code>Patience = 5</code> <code>optimizer=keras.optimizers.Adam</code>
<code>max_depth = [1, 6]</code> <code>n_estimators = 300</code> <code>max_features = {3, 5, 10, 20, 30, 50}</code>		<code>max_depth = [1, 2]</code> <code>n_estimators = [1,1000]</code> <code>learning_rate = [0.001,0.1]</code>	<code>max_depth = [1, 2]</code> <code>learning_rate = [0.001,0.1]</code> <code>n_estimators = [50,500]</code> <code>num_leaves = [10,100]</code> <code>min_child_samples = [1,50]</code> <code>subsample = [0.1,1]</code> <code>colsample_bytree = [0.1,1]</code>	<code>l1 = (1e-5 , 1e-3)</code> <code>learning_rate = {1e-3, 1e-2}</code> <code>batch_size = 10000</code> <code>epochs = 100</code> <code>Patience = 5</code> <code>optimizer=keras.optimizers.Adam</code> <code>activation = {relu, sigmoid}</code> <code>num_neurons[i] = (8, 256)</code>

Remark: Hyperparameters in the **first row** are chosen by *Shihao Gu (2020)* and hyperparameters in the **second row** are hyperparameters used in my program.


R-squared





Feature Importance

- Here we only display the top 10 important features for each model.



- Best Model: **NN3**
- Conicde with *Gu(2020)*
- Negative R concern

Negative R concern

Out of sample R squared:

$$R_{OOS}^2 = 1 - \frac{\sum (ret_{OOS} - \widehat{ret}_{OOS})^2}{\sum ret_{OOS}^2} = 1 - \frac{MSE \ Loss_{OOS}}{\sum ret_{OOS}^2}$$

Objective Function for Model in Training:

$$MSE \ Loss = \sum (ret_{IS} - \widehat{ret}_{IS})^2$$

- We try to minimize out of sample *MSE Loss* through minimizing insample *MSE Loss*. It's also worth to note that OLS estimator is the BLUE estimator. However, when we fixed all $\widehat{ret}_{OOS} = 0$, R_{OOS}^2 equals 0, which means that in some cases, the BLUE estimator performs worse than fixing every predicted value to 0.
- Simultaneous variation in R_{OOS}^2 performance within different model might implies our data has low ability in capture the volatility of stock market.
- From my perspective, negative R^2 might caused by some weaknesses of data.

Weaknesses of Data

Noise and Multi-collinearity

I use PCA to denoise and eliminate multi-collinearity exists between the variables.

```
#dimension-reduction method
pca = PCA(n_components='mle',svd_solver='full')
pca.fit(X_ptrain)
X_ptrain_pca = pca.transform(X_ptrain)
X_test_pca = pca.transform(X_test)
```

Remark: Another reason is to fasten the speed of training by reducing the number of variables.

Short-term trend

Add an indicator function to capture short term trend (rise/fall) per stock:

```
data['trend'] = data.apply(lambda x: 1 if data['ret'] - data['ret_l1'] > 0 else 0 if
data['ret'] - data['ret_l1'] == 0 else -1, axis=1) #avoid data leakage
```

Do contribute to R_{OOS}^2 performance when trend in stock market is stable.

Weaknesses of Data

Volatility of the stock market

- Our data show poor ability in capture the volatility of the stock market due to **leak of characteristics**. *Gu(2020)* uses 94 characteristics and we are only allowed to construct 63 of them.
- No **β Factor** in our data. Hard to capture stock's volatility in relation to the market.

Other Attempt

Training Neural Network with Pytorch using GPU

Activate GPU device:

```
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

Data Preparation:

```
# 1. convert data into tensor format
```

```
X_train = torch.tensor(X_train, dtype = torch.float)
```

```
y_train = torch.tensor(train_y, dtype = torch.float)
```

```
#2. transport data to device
```

```
X_train = X_train.to(device)
```

```
y_train = y_train.to(device)
```


Training Neural Network with Pytorch using GPU

Model Preparation:

```
#transport model to device  
net = AssetPricingNN(**best_params).to(device)
```

Detach data to numpy:

```
#transport data back to cpu  
predict_y = predict_y.detach().cpu().numpy()
```

```

class AssetPricingNN(nn.Module):
    def __init__(self, input_size, hidden_size_0, hidden_size_1, hidden_size_2,
output_size):
    super().__init__()
    self.emb0 = nn.Embedding(40000, 64)
    self.emb1 = nn.Embedding(40000, 64)
    self.emb2 = nn.Embedding(4000, 16)
    self.emb3 = nn.Embedding(4000, 16)
    self.fc1 = nn.Linear(input_size, hidden_size_0)
    self.fc2 = nn.Linear(hidden_size_0, hidden_size_1)
    self.fc3 = nn.Linear(hidden_size_1, hidden_size_2)
    self.predict = nn.Linear(hidden_size_2, output_size)
    self.emb = nn.Embedding(3, 2)

    def forward(self, x, cat0, cat1, cat2, cat3):
        cat0 = cat0.long()
        cat1 = cat1.long()
        cat2 = cat2.long()
        cat3 = cat3.long()

        cat0 = self.emb0(cat0)
        cat1 = self.emb1(cat1)
        cat2 = self.emb2(cat2)
        cat3 = self.emb3(cat3)

        x = torch.cat((x, cat0, cat1, cat2, cat3), dim = 1)

        x = torch.relu(self.fc1(x))
        x = torch.relu(self.fc2(x))
        x = torch.relu(self.fc3(x))
        x = self.predict(x)
        return x

```

Embedding

- In comparison to model that delete categorical variables directly, NN3 with embedding layers show no significant improvement in R_{OOS}^2 performance and run slower.

```

#Model architecture
class FactorAE(nn.Module):

    def __init__(self):
        super(FactorAE, self).__init__()
        #encoder architecture
        self.factor_output_layer = nn.Linear(58,3)
        self.batch1 = nn.BatchNorm2d(1,eps = 1e-5,affine = True)
        self.batch2 = nn.BatchNorm2d(1,eps = 1e-5,affine = True)
        self.beta_layer1 = nn.Linear(1,32)
        self.beta_layer2 = nn.Linear(32,16)
        self.beta_layer3 = nn.Linear(16,3)
        self.relu = nn.ReLU()

    def forward(self,data,y_return):

        OLS = []
        for i in range(X.shape[0]):
            x = X[i]
            x = x.reshape(1,-1)
            try:
                a = np.linalg.inv(np.dot(x.T,x))
            except:
                a = np.linalg.pinv(np.dot(x.T,x))
            b = np.dot(x.T,y[i])
            OLS.append(np.dot(a,b))

        OLS_tensor = OLS_tensor.squeeze(2)

        factor_output = self.factor_output_layer(OLS_tensor)
        #beta part
        beta = self.relu(self.batch1(self.beta_layer1(data.unsqueeze(1))))
        beta = self.relu(self.batch2(self.beta_layer2(beta)))
        beta = self.beta_layer3(beta).squeeze(1)
        reconstuct_return =
        torch.matmul(beta,factor_output.unsqueeze(2)).squeeze(2)

    return reconstuct_return

```

AutoEncoder

Pervious Model:

- use data before next month to generate parameters for model
- use data in **current month** to make prediction

AutoEncoder:

- use data before next month to generate parameters for model
- use **all historical data** to make prediction

Motivation: To capture **market volaitility** by model itself using historical data